

JHANVI SONI

Data Engineer | Databricks Platform Engineer | Analytics Engineer

Houston, TX • +1 (832) 246-6775 • jhanvisoni0@gmail.com • [linkedin.com/in/jhanvisonii/](https://www.linkedin.com/in/jhanvisonii/) • github.com/Jhanvisoni0

PROFESSIONAL SUMMARY

Data Engineer with 3.5+ years building production-grade ETL/ELT pipelines and analytics platforms across financial services (Morgan Stanley Capital International) and consulting (Techved). Specializes in PySpark, Medallion Architecture, and Delta Lake on Azure and Databricks, turning fragmented multi-source data into trusted dimensional models that cut reporting cycles from days to minutes. Seeking full-time Data, Analytics, or Azure Cloud Engineering roles with immediate availability. Experienced enabling and mentoring data teams on Databricks workflows, Delta Lake best practices, and production pipeline practices, helping users transition exploratory workflows into scalable, production-ready solutions.

TECHNICAL SKILLS & CERTIFICATIONS

Languages: Python, SQL, PySpark (Apache Spark), Java

Data Engineering & Orchestration: Azure Data Factory, Apache Airflow, Databricks Workflows, Auto Loader, Delta Live Tables, dbt, Apache Kafka

Cloud & Lakehouse: Azure (ADLS Gen2, Azure SQL, Synapse, Key Vault, Monitor, DevOps, Functions), Microsoft Fabric, Databricks, AWS

DevOps & IaC: Azure DevOps, CI/CD pipelines, Terraform, GitHub Actions, Git

Data Modeling & Architecture: Medallion Architecture (Bronze/Silver/Gold), Delta Lake, Star Schema, Dimensional Modeling, SCD Type 1/2

BI & Analytics Tools: Power BI, Tableau, MS SQL Server, Oracle, Advanced Excel, GitHub, Jupyter Notebooks

Analytical Methods: Regression, Time-Series Forecasting, A/B Testing, Hypothesis Testing, REST API / OpenAPI integration

Certifications: Azure Databricks Data Engineer Associate ([Credential](#)) • AWS Cloud Practitioner ([Credential](#)) • Databricks Fundamentals ([Credential](#))

EXPERIENCE

Data Engineer / Data Analyst | Morgan Stanley Capital International (MSCI) — Mumbai, India

Aug 2022 – July 2024

- Reduced average query execution time by **40%** through **partition pruning**, **predicate pushdown**, and **Z-ordering on high-cardinality** columns across **equity** and **Real Assets** reporting pipelines.
- Lifted reporting accuracy by **20%** by **building automated data quality frameworks** in **Python** monitoring **100+ feeds** covering **schema mismatch detection**, **null checks**, **duplicate alerts**, and **audit logging** catching issues before they reached end users.
- Migrated **20+ legacy SQL workflows** to **PySpark** on **Azure Databricks**; designed **star-schema fact/dimension tables** and **SCD Type 1 and Type 2** models that became the **single source of truth** for **quant**, **risk**, and **product teams**.
- Architected **ETL/ELT pipelines** on **Microsoft Azure (ADLS Gen2, Azure SQL, ADF)** to ingest and deliver large-scale **equities**, **fixed-income**, and **datasets**; provisioned infrastructure via **Terraform** and deployed pipeline changes through **Azure DevOps CI/CD** with automated validation gates.
- Built self-serve **Power BI** semantic models surfacing financial KPIs for stakeholder decision-making; mentored and **enabled** junior analysts on **Azure Databricks workflows**, **Delta Lake best practices**, and **production pipeline practices and platform adoption**.

Senior Data Analyst | Techved Consultancy — Mumbai, India

Jan 2021 – July 2022

- Built **SQL** and **Python** data pipelines, dimensional models, and **Power BI/Tableau** dashboards across **6+ client engagements** **troubleshooting requirements**, **enabling users**, and replacing manual reporting with automated, real-time operational visibility.
- Contributed to **\$300K+** in client expansion revenue by delivering **proof-of-concept** data models and dashboards; promoted from **Associate** to **Senior Associate** within **11 months**.
- Conducted deep-dive analysis using **Advanced Excel (pivot tables, XLOOKUP, reconciliation)** and **SQL**, validating data accuracy ahead of stakeholder-facing presentations.

EDUCATION

- M.S. in Management Information Systems — University of Houston Clear Lake, TX, USA
- B.A. in Psychology — University of Mumbai, MH, India

Aug 2024 – May 2026
July 2017 – May 2020

PROJECTS

Retail Banking Customer Analytics Platform - Azure · Databricks · PySpark · Delta Lake · ADF · Kafka · Power BI · Terraform

[View Project](#)

- Cut reporting cycle from **4 days to 15 minutes** by engineering a **Gold-layer customer-360 view** computing loan-to-deposit ratio and transaction velocity, automated via **Databricks Workflows** and deployed through a **Terraform-provisioned Azure** environment with **CI/CD pipelines**.
- Ingested **50K** transaction records and **5K** customer profiles from **3 source** systems into **ADLS Gen2** via **ADF pipelines** and **Databricks Auto Loader**; simulated real-time transaction event streaming using **Apache Kafka** topics feeding the **Bronze** ingestion layer.
- Secured credentials via **Azure Key Vault**, enforced table-level **RBAC** via **Unity Catalog**, delivered a 3-page executive **Power BI** dashboard via **DirectQuery**, and flagged anomalous accounts using a batch-scored **Isolation Forest** model tracked in **MLflow**.

Real Assets Investment Intelligence Platform - Databricks · PySpark · Delta Lake · yfinance & FRED APIs · MLflow · Unity Catalog

[View Project](#)

- Engineered an end-to-end investment analytics platform on **Azure Databricks** using **Medallion Architecture (Bronze/Silver/Gold)**, ingesting **22,635 rows** of live market data for **15 assets** via **yfinance** and **FRED APIs** into **Delta Lake tables** on **Unity Catalog Volumes**, with data partitioning, schema enforcement, and lineage tracking throughout.
- Built **PySpark** transformation pipelines computing portfolio KPIs daily returns, rolling volatility, Sharpe ratio, max drawdown, and S&P 500 benchmark correlation tracking a simulated **\$10K portfolio to \$18,645 (86% return)** over 6 years across **REITs**, infrastructure, and commodities.
- Joined **Fed Funds Rate** macro data (**FRED API**) with daily portfolio returns to reveal **REITs** returned **23.26%** annually in low-rate vs. - 8.17% in rate-transition periods; surfaced insights via a live **Databricks executive dashboard** with 90-day return forecasts.